

Assignment

Linear Regression: Real Estate Training Dataset

Prepare four Models where Price as Dependent Variable with each other sqft_living, Bedroom, Bathroom, floors as Independent Variables.

Model 1: Price as DV & sqft_living as IDV

Step 1: Load the dataset

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
lr_realestate = pd.read_excel("D:/AI_ML_Course/Day21/Linear Regression.xlsx")
```

```
lr_realestate.columns
```

```
Out[28]: Index(['price', 'sqft_living', 'bedrooms', 'bathrooms', 'floors'], dtype='object')
```

```
import sklearn
```

```
y = lr_realestate.price
```

```
X = lr_realestate[['sqft_living']]
```

Step 2: Split the records for training & testing

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 2)
```

```
X_train.shape
```

```
Out[35]: (17290, 1)
```

```
X_test.shape
```

```
Out[36]: (4323, 1)
```

```
y_train.shape
```

```
Out[37]: (17290,)
```

```
y_test.shape
```

```
Out[38]: (4323,)
```

Step 3: Model Building with sklearn

```
from sklearn.linear_model import LinearRegression
```

```
lin_reg = LinearRegression()
```

Step 4: Train the model

```
lin_reg.fit(X_train, y_train)
```

```
lin_reg.coef_
```

```
Out[42]: array([280.67382569])
```

```
lin_reg.intercept_
```

```
Out[43]: -42568.70358496299
```

Step 5: Visualize Training set result

```
plt.scatter(X_train, y_train, color = 'red')
```

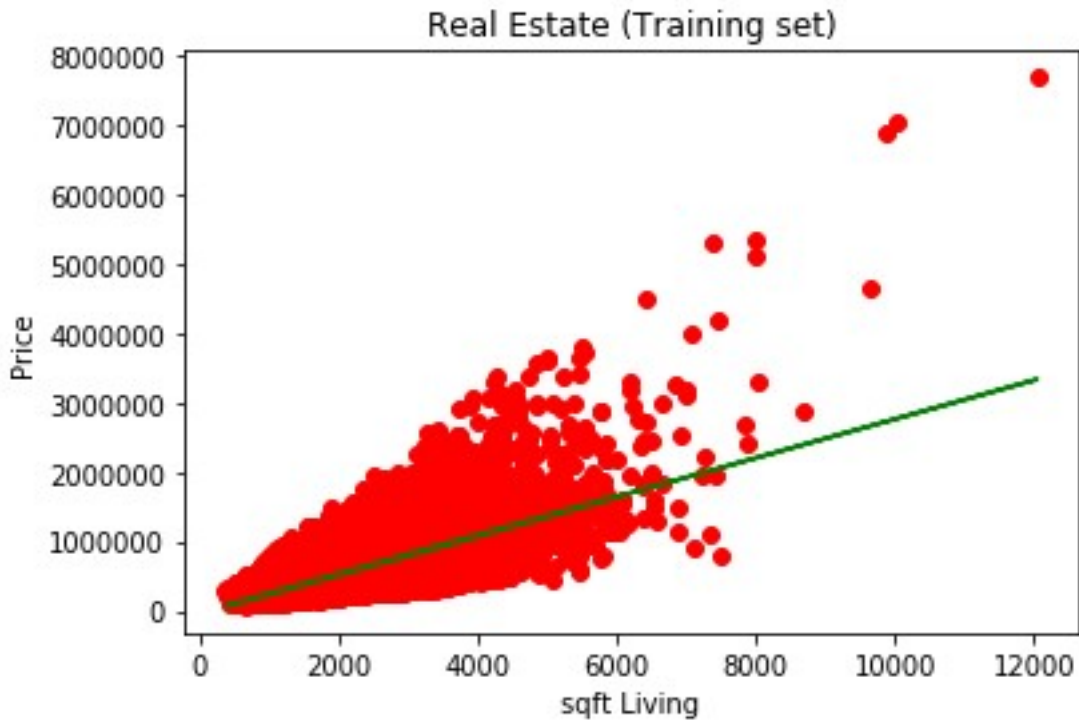
```
plt.plot(X_train, lin_reg.predict(X_train), color = 'green')
```

```
plt.title('Real Estate (Training set)')
```

```
plt.xlabel('sqft_living')
```

```
plt.ylabel('Price')
```

```
plt.show()
```



Step 6: Test the Model

```
ypred=lin_reg.predict(X_test)
```

```
print(ypred)
```

```
array([633855.21632509, 566493.49815977, 364408.34366382, ...,  
       701216.9344904 , 187583.83347987, 465450.9209118 ])
```

```
X_test.head()
```

```
Out[111]:
```

```
sqft_living  
6638    2410  
7366    2170  
3158    1450  
9117    4500  
3392     860
```

Step 7: Estimate the cost

```
from sklearn.metrics import mean_squared_error,r2_score  
  
RMSE=np.sqrt(mean_squared_error(y_test,ypred))  
  
r_square=r2_score(y_test,ypred)  
  
print('The R-Square value is...',r_square)
```

The R-Square value is... 0.5031163723285275

```
print('The RMSE value is.....',RMSE)
```

The RMSE value is..... 263380.00189817196

Step 8: How to predict for unseen value

```
unseen_pred=lin_reg.predict(np.array([[2410]]))  
  
print('The unseen for the given x is....',unseen_pred)
```

The unseen for the given x is.... [633855.21632509]

Model 2: Price as DV & Bedroom as IDV

Step 1: Load the dataset

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
lr_realestate = pd.read_excel("D:/AI_ML_Course/Day21/Linear Regression.xlsx")
```

```
lr_realestate.columns
```

```
Out[28]: Index(['price', 'sqft_living', 'bedrooms', 'bathrooms', 'floors'], dtype='object')
```

```
import sklearn
```

```
y = lr_realestate.price
```

```
X = lr_realestate[['bedrooms']]
```

Step 2: Split the records for training & testing

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 2)
```

```
X_train.shape
```

```
Out[35]: (17290, 1)
```

```
X_test.shape
```

```
Out[36]: (4323, 1)
```

```
y_train.shape
```

```
Out[37]: (17290,)
```

```
y_test.shape
```

```
Out[38]: (4323,)
```

Step 3: Model Building with sklearn

```
from sklearn.linear_model import LinearRegression
```

```
lin_reg = LinearRegression()
```

Step 4: Train the model

```
lin_reg.fit(X_train, y_train)
```

```
lin_reg.coef_
```

```
array([118660.62797869])
```

```
lin_reg.intercept_
```

```
139952.8759338616
```

Step 5: Visualize Training set result

```
plt.scatter(X_train, y_train, color = 'red')
```


```
plt.plot(X_train, lin_reg.predict(X_train), color = 'green')
```

```
plt.title('Real Estate (Training set)')
```

```
plt.xlabel('sqft_living')
```

```
plt.ylabel('Price')
```

```
plt.show()
```



Step 6: Test the Model

```
ypred=lin_reg.predict(X_test)
```

```
print(ypred)
```

```
array([614595.3878486 , 495934.75986992, 377274.13189123, ...,  
       614595.3878486 , 377274.13189123, 614595.3878486 ])
```

```
X_test.head()
```

```
Out[135]:
```

```
   bedrooms  
6638      4  
7366      3  
3158      2  
9117      5  
3392      2
```

Step 7: Estimate the cost

```
from sklearn.metrics import mean_squared_error,r2_score  
  
RMSE=np.sqrt(mean_squared_error(y_test,ypred))  
  
r_square=r2_score(y_test,ypred)  
  
print('The R-Square value is...',r_square)
```

The R-Square value is... 0.10886345250291585

```
print('The RMSE value is.....',RMSE)
```

The RMSE value is..... 352717.9654187645

Step 8: How to predict for unseen value

```
unseen_pred=lin_reg.predict(np.array([[3]]))  
  
print('The unseen for the given x is....',unseen_pred)
```

The unseen for the given x is.... [377274.13189123]

Model 3: Price as DV & Bathroom as IDV

Step 1: Load the dataset

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
lr_realestate = pd.read_excel("D:/AI_ML_Course/Day21/Linear Regression.xlsx")
```

```
lr_realestate.columns
```

```
Out[28]: Index(['price', 'sqft_living', 'bedrooms', 'bathrooms', 'floors'], dtype='object')
```

```
import sklearn
```

```
y = lr_realestate.price
```

```
X = lr_realestate[['bathrooms']]
```

Step 2: Split the records for training & testing

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 2)
```

```
X_train.shape
```

```
Out[35]: (17290, 1)
```

```
X_test.shape
```

```
Out[36]: (4323, 1)
```

```
y_train.shape
```

```
Out[37]: (17290,)
```

```
y_test.shape
```

```
Out[38]: (4323,)
```

Step 3: Model Building with sklearn

```
from sklearn.linear_model import LinearRegression
```

```
lin_reg = LinearRegression()
```

Step 4: Train the model

```
lin_reg.fit(X_train, y_train)
```

```
lin_reg.coef_
```

```
array([249143.95803858])
```

```
lin_reg.intercept_
```

```
13073.99575
```

Step 5: Visualize Training set result

```
plt.scatter(X_train, y_train, color = 'red')
```


```
plt.plot(X_train, lin_reg.predict(X_train), color = 'green')
```

```
plt.title('Real Estate (Training set)')
```

```
plt.xlabel('sqft_living')
```

```
plt.ylabel('Price')
```

```
plt.show()
```



Step 6: Test the Model

```
ypred=lin_reg.predict(X_test)
```

```
print(ypred)
```

```
array([573647.90133969, 386789.93281076, 262217.95379147, ...,  
       635933.89084934, 262217.95379147, 698219.88035898])
```

```
X_test.head()
```

```
Out[170]:
```

```
   bathrooms  
6638      2.25  
7366      1.50  
3158      1.00  
9117      3.25  
3392      1.00
```

Step 7: Estimate the cost

```
from sklearn.metrics import mean_squared_error,r2_score  
  
RMSE=np.sqrt(mean_squared_error(y_test,ypred))  
  
r_square=r2_score(y_test,ypred)  
  
print('The R-Square value is...',r_square)
```

The R-Square value is... 0.28122887124177365

```
print('The RMSE value is.....',RMSE)
```

The RMSE value is..... 316774.90190998075

Step 8: How to predict for unseen value

```
unseen_pred=lin_reg.predict(np.array([[2.25]]))  
  
print('The unseen for the given x is....',unseen_pred)
```

The unseen for the given x is.... [573647.90133969]

Model 4: Price as DV & Floors as IDV

Step 1: Load the dataset

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
lr_realestate = pd.read_excel("D:/AI_ML_Course/Day21/Linear Regression.xlsx")
```

```
lr_realestate.columns
```

```
Out[28]: Index(['price', 'sqft_living', 'bedrooms', 'bathrooms', 'floors'], dtype='object')
```

```
import sklearn
```

```
y = lr_realestate.price
```

```
X = lr_realestate[['floors']]
```

Step 2: Split the records for training & testing

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 2)
```

```
X_train.shape
```

```
Out[35]: (17290, 1)
```

```
X_test.shape
```

```
Out[36]: (4323, 1)
```

```
y_train.shape
```

```
Out[37]: (17290,)
```

```
y_test.shape
```

```
Out[38]: (4323,)
```

Step 3: Model Building with sklearn

```
from sklearn.linear_model import LinearRegression
```

```
lin_reg = LinearRegression()
```

Step 4: Train the model

```
lin_reg.fit(X_train, y_train)
```

```
lin_reg.coef_
```

```
array([171376.44562902])
```

```
lin_reg.intercept_
```

```
283309.93245028483
```

Step 5: Visualize Training set result

```
plt.scatter(X_train, y_train, color = 'red')
```

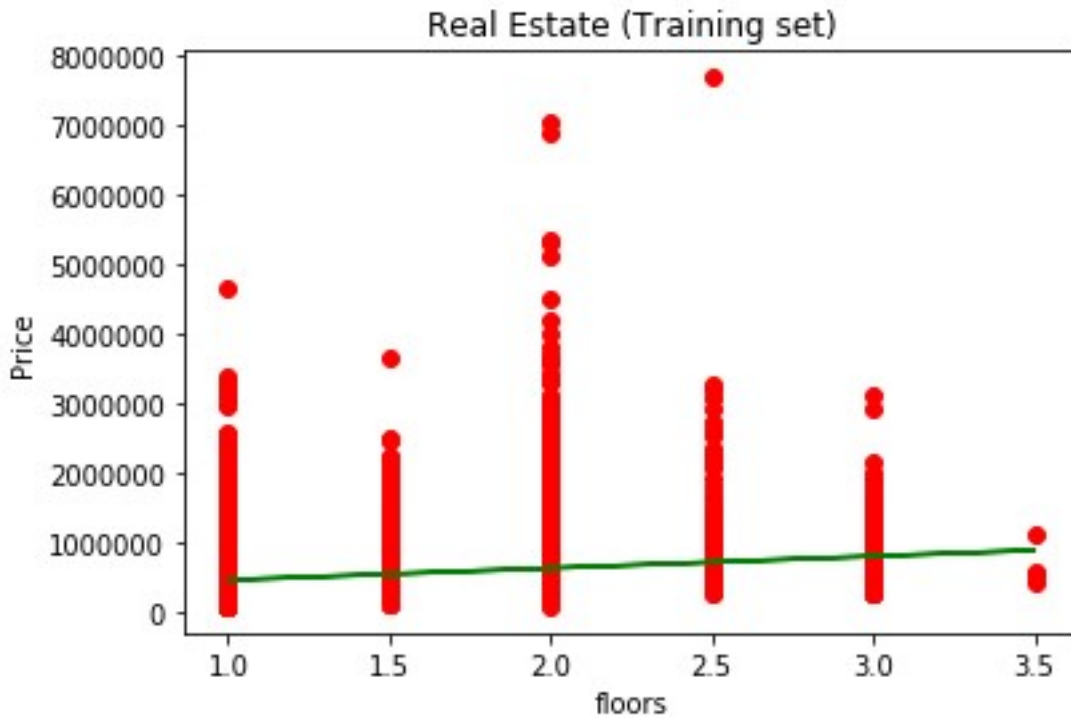
```
plt.plot(X_train, lin_reg.predict(X_train), color = 'green')
```

```
plt.title('Real Estate (Training set)')
```

```
plt.xlabel('floors')
```

```
plt.ylabel('Price')
```

```
plt.show()
```



Step 6: Test the Model

```
ypred=lin_reg.predict(X_test)
```

```
print(ypred)
```

```
array([540374.60089382, 454686.37807931, 454686.37807931, ...,
       626062.82370833, 454686.37807931, 540374.60089382])
```

```
X_test.head()
```

```
Out[197]:
```

```

      floors
6638    1.5
7366    1.0
3158    1.0
9117    2.0
3392    1.0
```

Step 7: Estimate the cost

```
from sklearn.metrics import mean_squared_error,r2_score  
  
RMSE=np.sqrt(mean_squared_error(y_test,ypred))  
  
r_square=r2_score(y_test,ypred)  
  
print('The R-Square value is...',r_square)
```

The R-Square value is... 0.0733487976687478

```
print('The RMSE value is.....',RMSE)
```

The RMSE value is..... 359677.77234107786

Step 8: How to predict for unseen value

```
unseen_pred=lin_reg.predict(np.array([[2.25]]))  
  
print('The unseen for the given x is....',unseen_pred)
```

The unseen for the given x is.... [626062.82370833]

Logistic Regression

Project 1: Bank Loan Modeling

Dependent Variable as “Personal Loan” and Independent Variable as “others”

Step 1: Load the Dataset

```
import pandas as pd

bankPL_dataset =
pd.read_excel("D:/AI_ML_Course/Day24/Bank_Personal_Loan_Modelling.xlsx",
sheet_name='Data')

bankPL_dataset.columns

# Index(['ID', 'Age', 'Experience', 'Income', 'ZIP Code', 'Family', 'CAvg',
#       'Education', 'Mortgage', 'Personal Loan', 'Securities Account',
#       'CD Account', 'Online', 'CreditCard'],
#       dtype='object')

bankPL_dataset = bankPL_dataset.drop(['ID', 'ZIP Code'], axis=1)

bankPL_dataset.columns

# Index(['Age', 'Experience', 'Income', 'Family', 'CAvg', 'Education',
#       'Mortgage', 'Personal Loan', 'Securities Account', 'CD Account',
#       'Online', 'CreditCard'],
#       dtype='object')
```

Step 2: Logistic Regression

```
import statsmodels.api as sm

Y = bankPL_dataset['Personal Loan']

X = bankPL_dataset[['Age', 'Experience', 'Income', 'Family', 'CAvg', 'Education', 'Mortgage',
'Securities Account', 'CD Account', 'Online', 'CreditCard']]
```

```

X1 = sm.add_constant(X)

BankKPL = sm.Logit(Y, X1)

result = BankKPL.fit()

# Optimization terminated successfully.

#      Current function value: 0.128435

#      Iterations 9

result.summary()

```

Logit Regression Results

```

=====
Dep. Variable:      Personal Loan          No. Observations:      5000
Model:              Logit                  Df Residuals:           4988
Method:             MLE                    Df Model:               11
Date:              Wed, 12 Aug 2020         Pseudo R-squ.:         0.5938
Time:              16:42:50                 Log-Likelihood:        -642.18
converged:          True                     LL-Null:               -1581.0
Covariance Type:    nonrobust                LLR p-value:           0.000
=====

```

	coef	std err	z	P> z	[0.025	0.975]

const	-12.1928	1.645	-7.411	0.000	-15.417	-8.968
Age	-0.0536	0.061	-0.874	0.382	-0.174	0.067
Experience	0.0638	0.061	1.046	0.295	-0.056	0.183
Income	0.0546	0.003	20.831	0.000	0.049	0.060
Family	0.6958	0.074	9.364	0.000	0.550	0.841

CCAvg	0.1240	0.040	3.127	0.002	0.046	0.202
Education	1.7362	0.115	15.088	0.000	1.511	1.962
Mortgage	0.0005	0.001	0.856	0.392	-0.001	0.002
Securities Account	-0.9368	0.286	-3.277	0.001	-1.497	-0.377
CD Account	3.8225	0.324	11.800	0.000	3.188	4.457
Online	-0.6752	0.157	-4.298	0.000	-0.983	-0.367
CreditCard	-1.1197	0.205	-5.462	0.000	-1.522	-0.718

=====

Step 3: Find the significant variables

With the above Logit Regression table, the variable where the P value is less than 0.05 is the significant variable.

So the significant variables are:

- Income
- Family
- CCAvg
- Education
- Securities Account
- CD Account
- Online
- CreditCard

Project 2: Attrition Rate Analysis

Dependent Variable as “Attrition” and Independent Variable as “others”

Step 1: Load the Dataset

```
import pandas as pd

attrition_dataset = pd.read_csv("D:/AI_ML_Course/Day24/general_data.csv")

attrition_dataset.columns

# Index(['Age', 'Attrition', 'BusinessTravel', 'Department', 'DistanceFromHome',
#        'Education', 'EducationField', 'EmployeeCount', 'EmployeeID', 'Gender',
#        'JobLevel', 'JobRole', 'MaritalStatus', 'MonthlyIncome',
#        'NumCompaniesWorked', 'Over18', 'PercentSalaryHike', 'StandardHours',
#        'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear',
#        'YearsAtCompany', 'YearsSinceLastPromotion', 'YearsWithCurrManager'],
#        dtype='object')

attrition_dataset = attrition_dataset.drop(['EmployeeCount',
'EmployeeID','Over18','StandardHours'], axis=1)

attrition_dataset.columns

# Index(['Age', 'Attrition', 'BusinessTravel', 'Department', 'DistanceFromHome',
#        'Education', 'EducationField', 'Gender', 'JobLevel', 'JobRole',
#        'MaritalStatus', 'MonthlyIncome', 'NumCompaniesWorked',
#        'PercentSalaryHike', 'StockOptionLevel', 'TotalWorkingYears',
#        'TrainingTimesLastYear', 'YearsAtCompany', 'YearsSinceLastPromotion',
#        'YearsWithCurrManager'],
#        dtype='object')
```

```

from sklearn import preprocessing

attrition_dataset = attrition_dataset.dropna()

attrition_dataset = attrition_dataset.drop_duplicates()

le = preprocessing.LabelEncoder()

attrition_dataset['Attrition'] = le.fit_transform(attrition_dataset['Attrition'])

attrition_dataset['BusinessTravel'] = le.fit_transform(attrition_dataset['BusinessTravel'])

attrition_dataset['Department'] = le.fit_transform(attrition_dataset['Department'])

attrition_dataset['EducationField'] = le.fit_transform(attrition_dataset['EducationField'])

attrition_dataset['Gender'] = le.fit_transform(attrition_dataset['Gender'])

attrition_dataset['JobRole'] = le.fit_transform(attrition_dataset['JobRole'])

attrition_dataset['MaritalStatus'] = le.fit_transform(attrition_dataset['MaritalStatus'])

```

Step 2: Logistic Regression

```

import statsmodels.api as sm

Y = attrition_dataset['Attrition']

X = attrition_dataset[['Age', 'BusinessTravel', 'Department', 'DistanceFromHome', 'Education',
'EducationField', 'Gender', 'JobLevel', 'JobRole', 'MaritalStatus', 'MonthlyIncome',
'NumCompaniesWorked', 'PercentSalaryHike', 'StockOptionLevel', 'TotalWorkingYears',
'TrainingTimesLastYear', 'YearsAtCompany', 'YearsSinceLastPromotion',
'YearsWithCurrManager']]

X1 = sm.add_constant(X)

Logistic_Attrition = sm.Logit(Y, X1)

result = Logistic_Attrition.fit()

# Optimization terminated successfully.

# Current function value: 0.392756

# Iterations 7

```

result.summary()

Logit Regression Results

```
=====
Dep. Variable:      Attrition                No. Observations:      1470
Model:              Logit                    Df Residuals:            1450
Method:             MLE                      Df Model:                19
Date:              Sun, 16 Aug 2020          Pseudo R-squ.:          0.1108
Time:              23:20:53                  Log-Likelihood:         -577.35
converged:          True                     LL-Null:                 -649.29
Covariance Type:    nonrobust                 LLR p-value:             3.295e-21
=====
```

```
=====
               coef  std err      z   P>|z|   [0.025   0.975]
-----
const          0.0650   0.717    0.091   0.928   -1.340    1.470
Age            -0.0306   0.012   -2.583   0.010   -0.054   -0.007
BusinessTravel -0.0166   0.113   -0.146   0.884   -0.239    0.206
Department     -0.2421   0.141   -1.720   0.085   -0.518    0.034
DistanceFromHome -0.0014   0.009   -0.145   0.884   -0.020    0.017
Education       -0.0625   0.074   -0.847   0.397   -0.207    0.082
EducationField  -0.0965   0.058   -1.669   0.095   -0.210    0.017
Gender          0.0869   0.155    0.560   0.576   -0.217    0.391
=====
```

JobLevel	-0.0249	0.069	-0.363	0.717	-0.159	0.110
JobRole	0.0378	0.031	1.219	0.223	-0.023	0.099
MaritalStatus	0.5885	0.109	5.379	0.000	0.374	0.803
MonthlyIncome	-1.868e-06	1.66e-06	-1.128	0.259	-5.11e-06	1.38e-06
NumCompaniesWorked	0.1184	0.032	3.729	0.000	0.056	0.181
PercentSalaryHike	0.0117	0.020	0.576	0.565	-0.028	0.052
StockOptionLevel	-0.0645	0.089	-0.721	0.471	-0.240	0.111
TotalWorkingYears	-0.0593	0.021	-2.856	0.004	-0.100	-0.019
TrainingTimesLastYear	-0.1465	0.061	-2.406	0.016	-0.266	-0.027
YearsAtCompany	0.0136	0.032	0.428	0.669	-0.049	0.076
YearsSinceLastPromotion	0.1323	0.035	3.732	0.000	0.063	0.202
YearsWithCurrManager	-0.1396	0.038	-3.642	0.000	-0.215	-0.064

=====

Step 3: Find the significant variables

With the above Logit Regression table, the variable where the P value is less than 0.05 is the significant variable.

So the significant variables are:

- Age
- Marital Status
- NumCompaniesWorked
- TotalWorkingYears
- TrainingTimeLastYear
- YearsSinceLastPromotion
- YearsWithCurrManager